

Deep Learning

Lab17: Deep RL

Pin-Yu Wang & Datalab

Outline

- From RL to Deep RL
- Deep Q -Network
 - Naïve Algorithm(TD)
 - Experience Replay
 - Delayed Target Network
 - Complete Algorithm
 - Implementation
- Assignment

Outline

- From RL to Deep RL
- Deep *Q*-Network
 - Naïve Algorithm(TD)
 - Experience Replay
 - Delayed Target Network
 - Complete Algorithm
 - Implementation
- Assignment

From RL to Deep RL

- (Tabular) RL
 - Q-learning:

$$Q^*(s, a) \leftarrow Q^*(s, a) + \eta[(R(s, a, s') + \gamma \max_{a'} Q^*(s', a')) - Q^*(s, a)]$$


- It requires a large table to store Q^* values in realistic environments with large state/action space.
 - Flappy bird: $O(10^5)$, Tetris: $O(10^{60})$, Automatic car: ???
- Hard to visit all (s, a) 's in limited training time.
- Agents must derive efficient representations of the environment from high-dimensional inputs and use these to **generalize past experience to new situations.**

Outline

- From RL to Deep RL
- Deep Q-Network
 - Naïve Algorithm(TD)
 - Experience Replay
 - Delayed Target Network
 - Complete Algorithm
 - Implementation
- Assignment

Deep Q-Network

- To learn a function $f_{Q^*}(s, a; \theta)$ that approximates $Q^*(s, a)$
 - Trained by a small number of samples.
 - Generalize to unseen states/actions.
 - Smaller θ to store.



Deep Q-Network

- Naïve Algorithm(TD)

1. Take action a from s using some exploration policy π' derived from f_{Q^*} (e.g. ϵ -greedy).
2. Observe s' and reward $R(s, a, s')$, and update θ using SGD:

$$\theta \leftarrow \theta - \eta \nabla_{\theta} C, \text{ where}$$

$$C(\theta) = [(R(s, a, s') + \gamma \max_{a'} f_{Q^*}(s', a'; \theta)) - f_{Q^*}(s, a; \theta)]^2$$

❖ Recall the Q-learning update formula:

$$Q^*(s, a) \leftarrow Q^*(s, a) + \eta [(R(s, a, s') + \gamma \max_{a'} Q^*(s', a')) - Q^*(s, a)]$$

Deep Q-Network

- However, the naïve TD algorithm diverges due to:
 1. Samples are correlated (violates i.i.d. assumption of training examples).
 2. Non-stationary target ($f_Q^*(s', a'; \theta)$ changes as θ is updated for current a).
- As a result, the Deep Q-Network applies two stabilization techniques to solve each problem respectively:
 1. Experience Replay
 2. Delayed Target Network

Outline

- From RL to Deep RL
- Deep Q-Network
 - Naïve Algorithm(TD)
 - Experience Replay
 - Delayed Target Network
 - Complete Algorithm
 - Implementation
- Assignment

Deep Q-Network

- Experience Replay
 - To remove the correlations in the observation sequence.
 - Use a replay memory \mathbb{D} to store recently seen transitions (s, a, r, s') 's.
 - Sample a mini-batch from \mathbb{D} and update θ .
 - The sample from the mini-batch is not a sequence now.

Outline

- From RL to Deep RL
- **Deep Q-Network**
 - Naïve Algorithm(TD)
 - Experience Replay
 - **Delayed Target Network**
 - Complete Algorithm
 - Implementation
- Assignment

Deep Q-Network

- Delayed Target Network

- To avoid chasing a moving target.
- Set the **target value** to the output of the network parameterized by **old θ^-** .
- Update $\theta^- \leftarrow \theta$ every K iterations.

❖ Update formula of naïve TD algorithm:

$$C(\theta) = [(R(s, a, s') + \gamma \max_{a'} f_{Q^*}(s', a'; \theta)) - f_{Q^*}(s, a; \theta)]^2$$

❖ Update formula after applying Delayed Target Network:

$$C(\theta) = [(R(s, a, s') + \gamma \max_{a'} f_{Q^*}(s', a'; \theta^-)) - f_{Q^*}(s, a; \theta)]^2$$

Outline

- From RL to Deep RL
- **Deep Q-Network**
 - Naïve Algorithm(TD)
 - Experience Replay
 - Delayed Target Network
 - **Complete Algorithm**
 - Implementation
- Assignment

Deep Q-Network

- Complete Algorithm

- Naïve algorithm(TD) + *Experience Replay* + *Delayed Target Network*
- Initialize θ arbitrarily and set $\theta^- = \theta$. Iterate until converge:
 1. Take action a from s using some exploration policy π' derived from f_{Q^*} (e.g. ε -greedy).
 2. Observe s' and reward $R(s, a, s')$, add (s, a, R, s') to \mathbb{D} .
 3. Sample a mini-batch of (s, a, R, s') 's from \mathbb{D} , do:

$$\theta \leftarrow \theta - \eta \nabla_{\theta} C, \text{ where}$$

$$C(\theta) = [(R(s, a, s') + \gamma \max_{a'} f_{Q^*}(s', a'; \theta^-)) - f_{Q^*}(s, a; \theta)]^2$$

4. Update $\theta^- \leftarrow \theta$ every K iterations.

Outline

- From RL to Deep RL
- **Deep Q-Network**
 - Naïve Algorithm(TD)
 - Experience Replay
 - Delayed Target Network
 - Complete Algorithm
 - **Implementation**
- Assignment

Outline

- From RL to Deep RL
- Deep *Q*-Network
 - Naïve Algorithm(TD)
 - Experience Replay
 - Delayed Target Network
 - Complete Algorithm
 - Implementation
- Assignment

Assignment – state-based DQN

- What you should do:
 - Change the input from stack of frames to game state(as Lab 16).
 - Change the network structure from CNNs to Dense layers.
 - Train the state-based DQN agent to play Flappy Bird.
- Evaluation metrics:
 - Code (Whether the implementation is correct) (50%).
 - The bird is able to fly through at least 1 pipes (50%).
- Requirements:
 - Upload the notebook and html file to google drive, and submit the link to iLMS.
 - Lab17_{student_id}.ipynb
 - Lab17_{student_id}.html
 - Deadline: 2020-12-31(Thur) 23:59.